

Problem Set 2

Zimo Shu

Problem 1 - Modified Random walk

- a. Implement the random walk in three versions.

Based on the rule, a single walk can only be: +1, -1, +10, -3.

Probabilities:

$$P(+1) = 0.5 \cdot 0.95 = 0.475, P(-1) = 0.5 \cdot 0.8 = 0.40, P(+10) = 0.5 \cdot 0.05 = 0.025, P(-3) = 0.5 \cdot 0.2 = 0.10.$$

(At first I was so not sure about the definition here so I reflected on the problem again. “If the +10 is chosen, it’s not +1 then +10, it is just +1.” \rightarrow If +1 is chosen, it could either be +1 or +10.)

```
steps <- c(10, 1, -1, -3)
probabilities <- c(0.025, 0.475, 0.40, 0.10)
```

Version 1: using a loop.

```
random_walk1 <- function(n) {
  start_pos <- 0 # start position at 0
  for (i in 1:n) {
    start_pos <- start_pos + sample(steps, 1, prob = probabilities, replace = TRUE)
  }
  start_pos
}
```

Version 2: using built-in R vectorized functions. (Using no loops.) (Hint: Does the order of steps matter?)

```
random_walk2 <- function(n) {  
  sum(sample(steps, n, replace = TRUE, prob = probabilities))  
}
```

Version 3: Implement the random walk using one of the “apply” functions.

```
random_walk3 <- function(n) {  
  sum(vapply(seq_len(n), function(i) sample(steps, 1, prob = probabilities, replace = TRUE),  
})
```

```
# Demonstrate that all versions work by running the following
```

```
random_walk1(10)
```

```
[1] 9
```

```
random_walk2(10)
```

```
[1] -8
```

```
random_walk3(10)
```

```
[1] 6
```

```
random_walk1(1000)
```

```
[1] 21
```

```
random_walk2(1000)
```

```
[1] 21
```

```
random_walk3(1000)
```

```
[1] -78
```

- b. Demonstrate that the three versions can give the same result. Show this for both $n=10$ and $n=1000$. (You will need to add a way to control the randomization.)

```
# set the same seed to control  
# for n = 10  
set.seed(506)  
random_walk1(10)
```

```
[1] -8
```

```
set.seed(506)  
random_walk2(10)
```

```
[1] -8
```

```
set.seed(506)  
random_walk3(10)
```

```
[1] -8
```

```
# for n = 10000  
set.seed(321)  
random_walk1(10000)
```

```
[1] 202
```

```
set.seed(321)  
random_walk2(10000)
```

```
[1] 202
```

```
set.seed(321)  
random_walk3(10000)
```

```
[1] 202
```

- c. Use the microbenchmark package to clearly demonstrate the speed of the implementations. Compare performance with a low input (1,000) and a large input (100,000). Discuss the results.

```
library(microbenchmark)
```

```
Warning: package 'microbenchmark' was built under R version 4.3.3
```

```
perform_func <- function(n, times = 50) {
  microbenchmark(
    loop  = random_walk1(n),
    apply = random_walk3(n),
    vect  = random_walk2(n),
    times = times
  )
}

bench_low <- perform_func(1000, 100) # low input
bench_large <- perform_func(100000, 100) # large input
bench_low
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval	cld
loop	4542.593	5301.276	5991.33397	5618.202	5931.4195	16765.147	100	a
apply	5473.722	5904.203	6594.10655	6119.053	6737.1790	11595.122	100	b
vect	25.815	27.414	30.15635	29.771	31.6105	54.782	100	c

```
bench_large
```

Unit: milliseconds

expr	min	lq	mean	median	uq	max	neval
loop	587.460210	628.980043	656.977458	650.0085	672.589971	816.275280	100
apply	615.900724	681.994970	711.825918	707.0139	729.341648	881.515425	100
vect	1.926122	2.233964	2.331749	2.3176	2.384133	2.860121	100

cld

a

b

c

Based on the outputs, in both low and large input cases, we could see that vectorized function (version 2) typically has the least running time, and fastest speed. Apply is the slowest. Loop is slightly faster than apply at these sizes.

We could also see the mean and median time of large inputs are way smaller than the mean and median time of low inputs, which is true for all three versions.

- d. What is the probability that the random walk ends at 0 if the number of steps is 10? 100? 1000? Defend your answers with evidence based upon a Monte Carlo simulation.

```
end_zero <- function(n, R = 1e6, seed = 123) {
  set.seed(seed)
  # Multinomial(n, p)
  X <- rmultinom(R, size = n, prob = probabilities)
  # final positions
  Final <- drop(steps %*% X)
  phat <- mean(Final == 0)      # Monte Carlo estimate
  se   <- sqrt(phat * (1 - phat) / R)
  ci   <- phat + qnorm(c(.025, .975)) * se
  list(p = phat, se = se, ci = ci)
}

res10 <- end_zero(10, R = 1e6, seed = 1)
res100 <- end_zero(100, R = 1e6, seed = 1)
res1000 <- end_zero(1000, R = 1e6, seed = 1)

res10
```

```
$p
[1] 0.132197
```

```
$se
[1] 0.0003387048
```

```
$ci
[1] 0.1315332 0.1328608
```

```
res100
```

```
$p
[1] 0.019751
```

```
$se
[1] 0.0001391434
```

```
$ci
[1] 0.01947828 0.02002372
```

```
res1000
```

```
$p  
[1] 0.0058  
  
$se  
[1] 7.593655e-05  
  
$ci  
[1] 0.005651167 0.005948833
```

We could tell that the probability of ending at 0 decreases with n. When n = 10, the probability is approximately 0.132. When n = 100, the probability is approximately 0.0197. When n = 10, the probability is approximately 0.0058.

Problem 2 - Mean of Mixture of Distributions

The number of cars passing an intersection is a classic example of a Poisson distribution. At a particular intersection, Poisson is an appropriate distribution most of the time, but during rush hours (hours of 8am and 5pm) the distribution is really normally distributed with a much higher mean.

Using a Monte Carlo simulation, estimate the average number of cars that pass an intersection per day under the following assumptions:

From midnight until 7 AM, the distribution of cars per hour is Poisson with mean 1. From 9am to 4pm, the distribution of cars per hour is Poisson with mean 8. From 6pm to 11pm, the distribution of cars per hour is Poisson with mean 12. During rush hours (8am and 5pm), the distribution of cars per hour is Normal with mean 60 and variance 12 Accomplish this without using any loops.

```
hour <- 0:23  
rush_hours <- c(8, 17) # 8am and 5pm  
normal_hours <- numeric(24)  
normal_hours[hour %in% 0:7] <- 1 # midnight until 7am  
normal_hours[hour %in% 9:16] <- 8 # 9am to 4pm  
normal_hours[hour %in% 18:23] <- 12 # 6pm to 11pm  
  
# Monte Carlo  
avg_cars <- function(R = 2e5, seed = 1) {  
  set.seed(seed)  
  pois_hours <- setdiff(hour, rush_hours)
```

```

# Poisson
X_pois <- matrix(rpois(length(pois_hours) * R, normal_hours[pois_hours + 1]), nrow = length(pois_hours))
# Rush hours
X_norm <- matrix(rnorm(2 * R, mean = 60, sd = sqrt(12)), nrow = 2)

# Daily totals
daily_totals <- colSums(X_pois) + colSums(X_norm)
# Monte Carlo estimate
estimation <- mean(daily_totals)
se <- sd(daily_totals) / sqrt(R)
ci <- estimation + qnorm(c(.025, .975)) * se
list(estimate = estimation, se = se, ci95 = ci)
}

# I got an error in using colsums: colSums(pmax(0, round(rbind(X_pois, X_norm)))) : 'x' must
avg_cars()

```

```
$estimate
[1] 263.9933
```

```
$se
[1] 0.02894918
```

```
$ci95
[1] 263.9366 264.0501
```

Hence, the average number of cars is approximately 264 (263.99).

Problem 3 - Linear Regression

```
youtube <- read.csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2022-01-18/youtube.csv')
```

a.

```
# Take a look at the table at first
head(youtube)
```

	year	brand
1	2018	Toyota
2	2020	Bud Light

3	2006	Bud Light				
4	2018	Hynudai				
5	2003	Bud Light				
6	2020	Toyota				
		superbowl_ads_dot_com_url				
1		https://superbowl-ads.com/good-odds-toyota/				
2		https://superbowl-ads.com/2020-bud-light-seltzer-inside-posts-brain/				
3		https://superbowl-ads.com/2006-bud-light-bear-attack/				
4		https://superbowl-ads.com/hope-detector-nfl-super-bowl-lii-hyundai/				
5		https://superbowl-ads.com/2003-bud-light-hermit-crab/				
6		https://superbowl-ads.com/2020-toyota-go-places-with-cobie-smulders/				
		youtube_url funny show_product_quickly				
1		https://www.youtube.com/watch?v=zeBZvwYQ-hA FALSE				
2		https://www.youtube.com/watch?v=nbbp0VW7z8w TRUE				
3		https://www.youtube.com/watch?v=yk0MQD5YgV8 TRUE				
4		https://www.youtube.com/watch?v=1NPccrGk77A FALSE				
5		https://www.youtube.com/watch?v=ovQYgnXHooY TRUE				
6		https://www.youtube.com/watch?v=f34Ji70u3nk TRUE				
		patriotic celebrity danger animals use_sex id kind				
1	FALSE	FALSE	FALSE	FALSE	zeBZvwYQ-hA	youtube#video
2	FALSE	TRUE	TRUE	FALSE	nbbp0VW7z8w	youtube#video
3	FALSE	FALSE	TRUE	TRUE	yk0MQD5YgV8	youtube#video
4	FALSE	FALSE	FALSE	FALSE	1NPccrGk77A	youtube#video
5	FALSE	FALSE	TRUE	TRUE	ovQYgnXHooY	youtube#video
6	FALSE	TRUE	TRUE	TRUE	f34Ji70u3nk	youtube#video
				etag view_count like_count dislike_count		
1		rn-ggKN1y38C10C3CNjNnUH9xUw	173929	1233	38	
2		1roDoK-SYqSpqYwKbYrMH0jEJQ4	47752	485	14	
3		0HiDfHTB3kilXfN8W0VTH0nwUIg	142310	129	15	
4		G9Dhby9Xe1UpnfcIrHmcnZYRCFI	198	2	0	
5		acsuYsFFQ_gHCC060Wus0tTgLjM	13741	20	3	
6		_UAuSOS1CytmE2NGd6u5pfaYhnA	23636	115	11	
		favorite_count comment_count published_at				
1		0	NA	2018-02-03T11:29:14Z		
2		0	14	2020-01-31T21:04:13Z		
3		0	9	2006-02-06T10:02:36Z		
4		0	0	2018-03-09T15:40:18Z		
5		0	2	2006-07-18T04:53:42Z		
6		0	13	2020-02-02T21:21:27Z		
		title				
1		Toyota Super Bowl Commercial 2018 Good Odds				
2		Bud Light: Post Malone #PostyStore Inside Post's Brain				
3		Super Bowl 2006: Bud Light "Save Yourself"				

```

4                               Hyundai / Hope Detector (2018)
5                               bud light pick up
6 Toyota Super Bowl Commercial 2020 Cobie Smulders Heroes

1
2 Bud Light, Post Malone "#PostyStore Inside Post's Brain"\n\nGarrick Sheldon (Copywriting, 2018)
3
4
5
6

          thumbnail      channel_title
1 https://i.ytimg.com/vi/zeBZvwYQ-hA/sddefault.jpg Funny Commercials
2 https://i.ytimg.com/vi/nbbp0VW7z8w/sddefault.jpg VCU Brandcenter
3                               <NA>      John Keehler
4                               <NA>      IATSE 490
5                               <NA>      jassymei
6 https://i.ytimg.com/vi/f34Ji70u3nk/sddefault.jpg Funny Commercials
category_id
1           1
2           27
3           17
4           22
5           24
6           1

# Remove any column that might uniquely identify a commercial
drop_cols <- c(
  "brand",
  "superbowl_ads_dot_com_url", "youtube_url", "thumbnail",
  "channel_title",
  "published_at",
  "id", "kind", "etag",
  "title", "description"
)

youtube_cleaned <- youtube[ , !(names(youtube) %in% drop_cols)]

# report dimensions
dim(youtube_cleaned)

```

[1] 247 14

The dimensional now is as above.

b.

```
vars <- c("view_count", "like_count", "dislike_count", "favorite_count", "comment_count")
# Examine the distribution
summ <- sapply(youtube[vars], \((x) c(
  n      = sum(!is.na(x)),
  min    = min(x, na.rm=TRUE),
  q25    = unname(quantile(x, .25, na.rm=TRUE)),
  median = median(x, na.rm=TRUE),
  mean   = mean(x, na.rm=TRUE),
  q75    = unname(quantile(x, .75, na.rm=TRUE)),
  max    = max(x, na.rm=TRUE),
  sd     = sd(x, na.rm=TRUE),
  p_zeros= mean(x == 0, na.rm=TRUE)
))
round(t(summ), 3)
```

	n	min	q25	median	mean	q75	max	sd
view_count	231	10	6431	41379	1407556.459	170015.50	176373378	11971111.010
like_count	225	0	19	130	4146.031	527.00	275362	23920.403
dislike_count	225	0	1	7	833.538	24.00	92990	6948.522
favorite_count	231	0	0	0	0.000	0.00	0	0.000
comment_count	222	0	1	10	188.640	50.75	9190	986.457
	p_zeros							
view_count		0.000						
like_count		0.040						
dislike_count		0.209						
favorite_count		1.000						
comment_count		0.185						

From the outputs, we could see:

view_count: (ii) right-skewed, should be transformed
like_count: right-skewed, outliers, should be transformed
dislike_count: skewed, outliers, should be transformed
favorite_count: sd is zero (variance is zero), **not appropriate**
comment_count: right-skewed, should be transformed.

```
# Transformation
youtube$view_count_log      <- log1p(youtube$view_count)
youtube$like_count_log       <- log1p(youtube$like_count)
youtube$dislike_count_log   <- log1p(youtube$dislike_count)
youtube$comment_count_log    <- log1p(youtube$comment_count)
```

- c. For each variable in part b. that are appropriate, fit a linear regression model predicting them based upon each of the seven binary flags for characteristics of the ads, such as whether it is funny. Control for year as a continuous covariate.

Discuss the results. Identify the direction of any statistically significant results.

```
# Seven binary flags
flags <- c("funny", "show_product_quickly", "patriotic",
          "celebrity", "danger", "animals", "use_sex")
flags <- flags[flags %in% names(youtube)]
```

```
# Fit models and print outputs
fit_models <- function(y) {
  formulas <- as.formula(paste(y, "~ year +", paste(flags, collapse = " + ")))
  dat <- youtube[, c(y, "year", flags)]
  dat <- dat[complete.cases(dat), ]
  mod <- lm(formulas, data = dat)
  cat("\n=====\n")
  cat("Outcome:", y, "\n")
  cat("=====\n")
  print(summary(mod)$coefficients)
  invisible(mod)
}

mods <- lapply(c("view_count_log", "like_count_log", "dislike_count_log", "comment_count_log"),
```

```
=====
Outcome: view_count_log
=====

              Estimate Std. Error   t value Pr(>|t|) 
(Intercept) -31.55015804 71.00537527 -0.4443348 0.6572334
year          0.02053399  0.03530599  0.5816007 0.5614258
funnyTRUE     0.56492445  0.46702107  1.2096338 0.2277060
show_product_quicklyTRUE 0.21088918  0.40530215  0.5203258 0.6033550
patrioticTRUE 0.50699051  0.53811043  0.9421681 0.3471307
```

celebrityTRUE	0.03547862	0.42228051	0.0840167	0.9331189
dangerTRUE	0.63131085	0.41812403	1.5098650	0.1324998
animalsTRUE	-0.31001838	0.39347937	-0.7878898	0.4316016
use_sexTRUE	-0.38670726	0.44781782	-0.8635370	0.3887743

=====

Outcome: like_count_log

=====

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-150.51356902	63.45723424	-2.3718898	0.01857564
year	0.07684959	0.03155325	2.4355517	0.01567989
funnyTRUE	0.47476026	0.41816138	1.1353518	0.25748620
show_product_quicklyTRUE	0.20017362	0.36390920	0.5500647	0.58284333
patrioticTRUE	0.80688663	0.49790889	1.6205507	0.10657295
celebrityTRUE	0.41283478	0.38212296	1.0803716	0.28118156
dangerTRUE	0.63894992	0.37350269	1.7106970	0.08857278
animalsTRUE	-0.05943769	0.35298000	-0.1683883	0.86643540
use_sexTRUE	-0.42951949	0.40063789	-1.0720890	0.28487641

=====

Outcome: dislike_count_log

=====

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-183.06812777	53.34768465	-3.4316040	0.0007189054
year	0.09207348	0.02652642	3.4710107	0.0006260350
funnyTRUE	0.25949148	0.35154292	0.7381502	0.4612243850
show_product_quicklyTRUE	0.27511424	0.30593381	0.8992607	0.3695152223
patrioticTRUE	0.81407206	0.41858564	1.9448160	0.0530952460
celebrityTRUE	-0.20214335	0.32124588	-0.6292481	0.5298516125
dangerTRUE	0.22184469	0.31399893	0.7065142	0.4806298632
animalsTRUE	-0.21191815	0.29674576	-0.7141405	0.4759113338
use_sexTRUE	-0.32980132	0.33681115	-0.9791877	0.3285826970

=====

Outcome: comment_count_log

=====

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-99.09835037	52.92351239	-1.8724825	0.06251009
year	0.05033845	0.02631887	1.9126376	0.05713584
funnyTRUE	0.21954414	0.34527939	0.6358449	0.52556010
show_product_quicklyTRUE	0.40938620	0.30229342	1.3542676	0.17708596
patrioticTRUE	0.66698069	0.39901933	1.6715498	0.09608156
celebrityTRUE	0.29767273	0.31541291	0.9437557	0.34636374

```

dangerTRUE          0.17783936  0.31068999  0.5724013  0.56765378
animalsTRUE         -0.26802431  0.29346559 -0.9133074  0.36211349
use_sexTRUE         -0.39323227  0.33162665 -1.1857680  0.23703504

```

From the outputs, we could see that for view_count_log, all $p > 0.13$, so no predictors are significant. For like_count_log, year is positive and significant. All flags are not significant. For dislike_count_log, year is positive and significant. Flag patriotic is very close to 0.05, and the direction is positive. All others are not significant. For comment_count_log, no predictors are really significant. Flag year is very close to 0.05 though.

Hence, after log-transforming, there is no sufficient evidence that shows that the seven flags (funny, show product quickly, patriotic, celebrity, danger, animals, use sex) are statistically associated with views, likes, dislikes, or comments at the 0.05 level. (We could see for like_count_log and dislike_count_log, the variable year is positive and significant, which indicates that year is significantly positive for likes and dislikes.)

- d. Consider only the outcome of view counts. Calculate $\hat{\beta}$ manually (without using lm) by first creating a proper design matrix, then using matrix algebra to estimate β . Confirm that you get the same result as lm did in part c.

```

fml_viewcount <- as.formula(
  paste("view_count_log ~ year +", paste(flags, collapse = " + "))
)

vars_new <- c("view_count_log", "year", flags)
dat <- youtube[complete.cases(youtube[, vars_new]), vars_new]

# design matrix X and response y
X <- model.matrix(fml_viewcount, data = dat)
y <- dat$view_count_log

# Use normal equations
XtX <- t(X) %*% X
Xty <- t(X) %*% y
beta_hat <- solve(XtX, Xty)

# Use lm to get the value
fit <- lm(fml_viewcount, data = dat)
coef_lm <- coefficients(fit)

print(beta_hat)

```

```
[,1]
(Intercept) -31.55015804
year          0.02053399
funnyTRUE     0.56492445
show_product_quicklyTRUE 0.21088918
patrioticTRUE 0.50699051
celebrityTRUE 0.03547862
dangerTRUE    0.63131085
animalsTRUE   -0.31001838
use_sexTRUE   -0.38670726
```

```
print(coef_lm)
```

(Intercept)	year	funnyTRUE
-31.55015804	0.02053399	0.56492445
show_product_quicklyTRUE	patrioticTRUE	celebrityTRUE
0.21088918	0.50699051	0.03547862
dangerTRUE	animalsTRUE	use_sexTRUE
0.63131085	-0.31001838	-0.38670726

```
all.equal(as.vector(beta_hat), as.vector(coef_lm))
```

```
[1] TRUE
```

Hence, I got the same result as lm did. The calculated result would be:

$$\hat{\beta} \approx (-31.55, 0.02, 0.56, 0.21, 0.51, 0.04, 0.63, -0.31, -0.38).$$